

Big Data in Economics

Lecture 7: Webscraping: (1) Server-side and CSS

Grant R. McDermott

University of Oregon | EC 607

Contents

Software requirements	1
Webscraping basics	1
Webscraping with rvest (server-side)	2
Application: Mens 100 meters (Wikipedia)	3
Summary	11
Further resources and exercises	12

Software requirements

External software

Today we'll be using SelectorGadget, which is a Chrome extension that makes it easy to discover CSS selectors. (Install the extension directly here.) Please note that SelectorGadget is only available for Chrome. If you prefer using Firefox, then you can try ScrapeMate.

R packages

- New: **rvest**, **janitor**
- Already used: **tidyverse**, **lubridate**, **hrbrthemes**

Recall that **rvest** was automatically installed with the rest of the tidyverse. Still, here is a convenient way to install (if necessary) and load all of the above packages.

```
## Load and install the packages that we'll be using today
if (!require("pacman")) install.packages("pacman")
pacman::p_load(tidyverse, rvest, lubridate, janitor, hrbrthemes)
## My preferred ggplot2 plotting theme (optional)
theme_set(hrbrthemes::theme_ipsum())
```

Tip: If you can get an error about missing fonts whilst following along with this lecture, that's probably because you don't have Arial Narrow — required by the `hrbrthemes::theme_ipsum()` **ggplot2** theme that I'm using here — installed on your system. You can resolve this by downloading the font and adding it to your font book (Google it), or by switching to a different theme (e.g. `theme_set(theme_minimal())`).

Webscraping basics

The next two lectures are about getting data, or “content”, off the web and onto our computers. We're all used to seeing this content in our browsers (Chrome, Firefox, etc.). So we know that it must exist somewhere. However, it's important to realise that there are actually two ways that web content gets rendered in a browser:

1. Server-side
2. Client side

You can read here for more details (including example scripts), but for our purposes the essential features are as follows:

1. Server-side

- The scripts that “build” the website are not run on our computer, but rather on a host server that sends down all of the HTML code.
 - E.g. Wikipedia tables are already populated with all of the information — numbers, dates, etc. — that we see in our browser.
- In other words, the information that we see in our browser has already been processed by the host server.
- You can think of this information being embedded directly in the webpage’s HTML.
- **Webscraping challenges:** Finding the correct CSS (or Xpath) “selectors”. Iterating through dynamic webpages (e.g. “Next page” and “Show More” tabs).
- **Key concepts:** CSS, Xpath, HTML

2. Client-side

- The website contains an empty template of HTML and CSS.
 - E.g. It might contain a “skeleton” table without any values.
- However, when we actually visit the page URL, our browser sends a *request* to the host server.
- If everything is okay (e.g. our request is valid), then the server sends a *response* script, which our browser executes and uses to populate the HTML template with the specific information that we want.
- **Webscraping challenges:** Finding the “API endpoints” can be tricky, since these are sometimes hidden from view.
- **Key concepts:** APIs, API endpoints

Over the next two lectures, we’ll go over the main differences between the two approaches and cover the implications for any webscraping activity. I want to forewarn you that webscraping typically involves a fair bit of detective work. You will often have to adjust your steps according to the type of data you want, and the steps that worked on one website may not work on another. (Or even work on the same website a few months later). All this is to say that *webscraping involves as much art as it does science*.

The good news is that both server-side and client-side websites allow for webscraping.¹ If you can see it in your browser, you can scrape it.

Caveat: Ethical and legal considerations

The previous sentence elides some important ethical considerations. Just because you *can* scrape it, doesn’t mean you *should*. Now, I first have to tell you that this paragraph used to contain a warning about the legal restrictions pertaining to webscraping activity. I’ve decided to drop those in the wake of the landmark *hiQ Labs vs LinkedIn* court ruling. (Short version: It is currently legal to scrape data from the web using automated tools, as long as the data are publicly available.) However, it’s still important to realise that the tools we’ll be using over these next two lectures are very powerful. A computer can process commands much, much faster than we can ever type them up manually. It’s pretty easy to write up a function or program that can overwhelm a host server or application through the sheer weight of requests. Or, just as likely, the host server has built-in safeguards that will block you in case of a suspected malicious attack. We’ll return to the “be nice” mantra at the end of this lecture, as well as in the next lecture.

Webscraping with rvest (server-side)

The primary R package that we’ll be using today is **rvest** (link), a simple webscraping library inspired by Python’s **Beautiful Soup** (link), but with extra tidyverse functionality. **rvest** is designed to work with web pages that are built server-side and thus requires knowledge of the relevant CSS selectors... Which means I should probably tell you what those are.

¹As we’ll see during the next lecture, scraping a website or application that is built on a client-side (i.e. API) framework is often easier; particularly when it comes to downloading information *en masse*.

CSS

CSS (i.e. “Cascading Style Sheets”) is a computer language for specifying the appearance of HTML documents like web pages. It does this by providing a set of display rules, the main features of which are:

1. *Properties*. CSS properties are the “how” of the display rules. These are things like which font family, styles and colours to use, page width, etc.
2. *Selectors*. CSS selectors are the “what” of the display rules. They identify which properties should be applied to which elements. E.g. Text elements that are selected as “.h1” (i.e. top line headers) will inherit larger font sizes and are usually displayed more prominently than text elements selected as “.h2” (i.e. sub-headers), etc.

While this provides some useful background, the key point for us is as follows: If you can identify the CSS selector(s) pertaining to your content of interest — e.g. a table or paragraph of text in a webpage — then you can isolate it from the rest of the content that you don’t want. And once you have isolated it, then you can scrape it...

But how do we identify the right CSS selectors in the first place? Well, I’m going to show you two ways: (i) Using the SelectorGadget browser extension that you should have installed already, and (ii) Using your browser’s built-in inspection tools. Let’s proceed with a real-life application.

Application: Mens 100 meters (Wikipedia)

Say that we want to scrape the Wikipedia page on the **Men’s 100 metres world record progression**.

First, open up this page in your browser and take a look at its structure. - What type of objects does it contain? - How many tables does it have? - Do these tables all share the same columns? - What about row- and columns-spans? - Etc.

Once you’ve familiarized yourself with this structure, read the whole page into R using the `rvest::read_html()` function.

```
# library(rvest) ## Already loaded

m100 <- read_html("http://en.wikipedia.org/wiki/Men%27s_100_metres_world_record_progression")
m100

## {html_document}
## <html class="client-nojs" lang="en" dir="ltr">
## [1] <head>\n<meta http-equiv="Content-Type" content="text/html; charset=UTF-8 ...
## [2] <body class="mediawiki ltr sitedir-ltr mw-hide-empty-elt ns-0 ns-subject ...
```

As you can see, this is an XML document² that contains *everything* needed to render the Wikipedia page. It’s kind of like viewing someone’s entire dissertation (table of contents, introduction, etc.) when all we want are the data from some tables in their paper.

Table 1: Pre-IAAF (1881–1912)

Let’s try to isolate the first table on the page, which documents the unofficial progression before the IAAF. As per the *rvest* vignette, we can use `rvest::html_nodes()` to isolate and extract this table from the rest of the HTML document by providing the relevant CSS selector. We should then be able to convert it into a data frame using `rvest::html_table()`. I also recommend using the `fill=TRUE` option here, because otherwise we’ll run into formatting problems due to row spans in the Wiki table.

I’ll start by using the SelectorGadget to identify the relevant CSS selector. I’ll walk through this in class, but watch the video at the link if you’re reading this after the fact and unsure of what to do.³ In this case, I get “div+.wikitable :nth-child(1)”, so let’s check if that works.

```
m100 %>%
  html_nodes("div+.wikitable :nth-child(1)") %>%
  html_table(fill=TRUE)
```

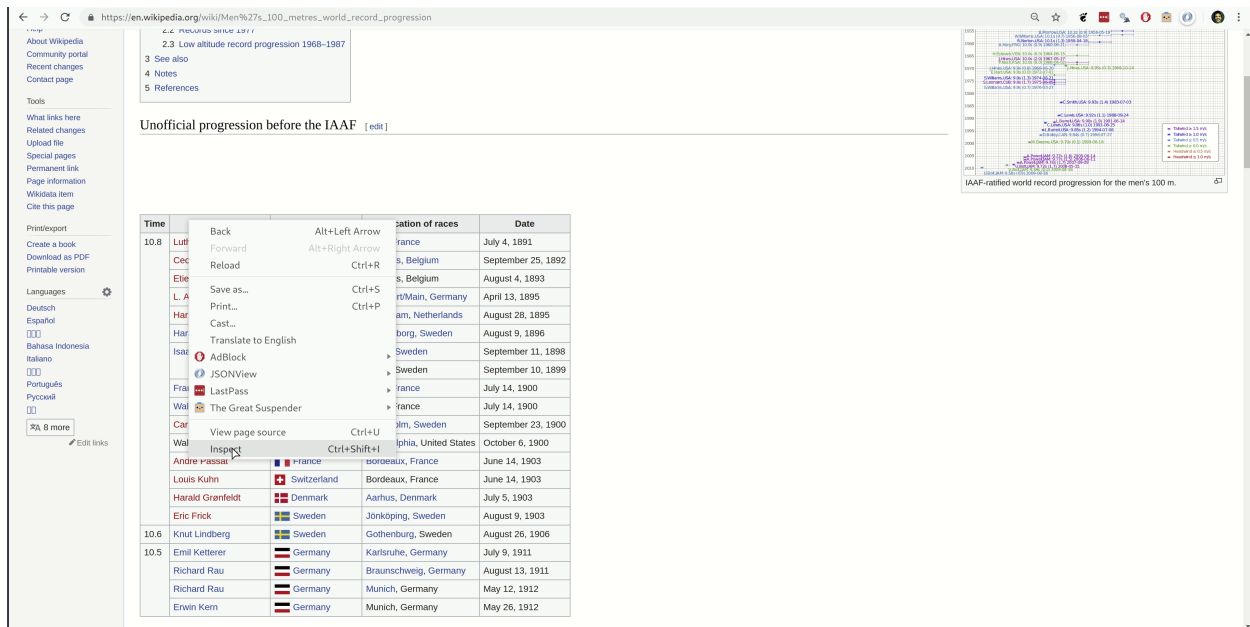
²XML stands for Extensible Markup Language and is one of the primary languages used for encoding and formatting web pages.

³The *rvest* package has a nice, short vignette with SelectorGadget screenshots if you’d prefer.

```
## Error in html_table.xml_node(X[[i]], ...): html_name(x) == "table" is not TRUE
```

Uh-oh! It seems that we immediately run into an error. I won't go into details here, but we have to be cautious with SelectorGadget sometimes. It's a great tool and usually works perfectly. However, occasionally what looks like the right selection (i.e. the highlighted stuff in yellow) is not exactly what we're looking for. I deliberately chose this Wikipedia 100m example because I wanted to showcase this potential pitfall. Again: Webscraping is as much art as it is science.

Fortunately, there's a more precise way of determining the right selectors using the "inspect web element" feature that is available in all modern browsers. In this case, I'm going to use Google Chrome (**Ctrl+Shift+I**, or right-click and choose "Inspect"). I proceed by scrolling over the source elements until Chrome highlights the table of interest. Then right-click again and choose **Copy -> Copy selector**. Here's a GIF animation of these steps:



Using this method, I get "#mw-content-text > div > table:nth-child(8)". Let's see whether it works this time. Again, I'll be using the `rvest::html_table(fill=TRUE)` function to coerce the resulting table into a data frame.

```
m100 %>%
  html_nodes("#mw-content-text > div > table:nth-child(8)") %>%
  html_table(fill=TRUE)
```

```
## [[1]]
##      Time      Athlete      Nationality      Location of races
## 1  10.8      Luther Cary  United States      Paris, France
## 2  10.8      Cecil Lee   United Kingdom      Brussels, Belgium
## 3  10.8      Étienne De Ré      Belgium      Brussels, Belgium
## 4  10.8      L. Atcherley  United Kingdom      Frankfurt/Main, Germany
## 5  10.8      Harry Beaton  United Kingdom      Rotterdam, Netherlands
## 6  10.8      Harald Anderson-Arbin      Sweden      Helsingborg, Sweden
## 7  10.8      Isaac Westergren      Sweden      Gävle, Sweden
## 8  10.8      10.8          Sweden      Gävle, Sweden
## 9  10.8      Frank Jarvis   United States      Paris, France
## 10 10.8      Walter Tewksbury  United States      Paris, France
## 11 10.8      Carl Ljung      Sweden      Stockholm, Sweden
## 12 10.8      Walter Tewksbury  United States      Philadelphia, United States
## 13 10.8      André Passat      France      Bordeaux, France
## 14 10.8      Louis Kuhn      Switzerland      Bordeaux, France
## 15 10.8      Harald Grønfeldt  Denmark      Aarhus, Denmark
```

```
## 16 10.8          Eric Frick          Sweden          Jönköping, Sweden
## 17 10.6          Knut Lindberg       Sweden          Gothenburg, Sweden
## 18 10.5          Emil Ketterer       Germany         Karlsruhe, Germany
## 19 10.5          Richard Rau         Germany         Braunschweig, Germany
## 20 10.5          Richard Rau         Germany          Munich, Germany
## 21 10.5          Erwin Kern          Germany          Munich, Germany
##                               Date
## 1              July 4, 1891
## 2 September 25, 1892
## 3              August 4, 1893
## 4              April 13, 1895
## 5              August 28, 1895
## 6              August 9, 1896
## 7 September 11, 1898
## 8 September 10, 1899
## 9              July 14, 1900
## 10             July 14, 1900
## 11 September 23, 1900
## 12             October 6, 1900
## 13             June 14, 1903
## 14             June 14, 1903
## 15             July 5, 1903
## 16             August 9, 1903
## 17             August 26, 1906
## 18             July 9, 1911
## 19             August 13, 1911
## 20             May 12, 1912
## 21             May 26, 1912
```

Great, it worked! Let's assign it to an object that we'll call `pre_iaaf` and then check its class.

```
pre_iaaf <-
  m100 %>%
    html_nodes("#mw-content-text > div > table:nth-child(8)") %>%
    html_table(fill=TRUE)
class(pre_iaaf)
```

```
## [1] "list"
```

Hmmm... It turns out this is actually a list, so let's *really* convert it to a data frame. You can do this in multiple ways. I'm going to use the `dplyr::bind_rows()` function, which is great for coercing (multiple) lists into a data frame.⁴

```
## Convert list to data_frame
# pre_iaaf <- pre_iaaf[[1]] ## Would also work

# library(tidyverse) ## Already loaded

pre_iaaf <-
  pre_iaaf %>%
    bind_rows() %>%
    as_tibble()
pre_iaaf
```

```
## # A tibble: 21 x 5
##   Time Athlete          Nationality `Location of races` Date
```

⁴We'll see more examples of this once we get to the programming section of the course.

```
##      <dbl> <chr>                <chr>                <chr>                <chr>
## 1  10.8 Luther Cary             United States Paris, France       July 4, 1891
## 2  10.8 Cecil Lee               United Kingdom Brussels, Belgium    September 25, ~
## 3  10.8 Étienne De Ré          Belgium        Brussels, Belgium    August 4, 1893
## 4  10.8 L. Atcherley           United Kingdom Frankfurt/Main, Germ~ April 13, 1895
## 5  10.8 Harry Beaton           United Kingdom Rotterdam, Netherlan~ August 28, 1895
## 6  10.8 Harald Anderson-A~ Sweden         Helsingborg, Sweden August 9, 1896
## 7  10.8 Isaac Westergren       Sweden         Gävle, Sweden        September 11, ~
## 8  10.8 10.8                   Sweden         Gävle, Sweden        September 10, ~
## 9  10.8 Frank Jarvis           United States  Paris, France        July 14, 1900
## 10 10.8 Walter Tewksbury       United States  Paris, France        July 14, 1900
## # ... with 11 more rows
```

Let's fix the column names to get rid of spaces, etc. I'm going to use the `janitor::clean_names()` function, which is expressly built for the purpose of cleaning object names. (Q: How else could we have done this?)

```
# library(janitor) ## Already loaded
```

```
pre_iaaf <-
  pre_iaaf %>%
    clean_names()
pre_iaaf
```

```
## # A tibble: 21 x 5
##   time athlete      nationality location_of_races    date
##   <dbl> <chr>          <chr>          <chr>          <chr>
## 1  10.8 Luther Cary   United States Paris, France    July 4, 1891
## 2  10.8 Cecil Lee    United Kingdom Brussels, Belgium September 25, ~
## 3  10.8 Étienne De Ré Belgium        Brussels, Belgium August 4, 1893
## 4  10.8 L. Atcherley United Kingdom Frankfurt/Main, Germ~ April 13, 1895
## 5  10.8 Harry Beaton United Kingdom Rotterdam, Netherlan~ August 28, 1895
## 6  10.8 Harald Anderson-A~ Sweden         Helsingborg, Sweden August 9, 1896
## 7  10.8 Isaac Westergren Sweden         Gävle, Sweden    September 11, ~
## 8  10.8 10.8         Sweden         Gävle, Sweden    September 10, ~
## 9  10.8 Frank Jarvis United States  Paris, France    July 14, 1900
## 10 10.8 Walter Tewksbury United States  Paris, France    July 14, 1900
## # ... with 11 more rows
```

Hmmm. There is a slight misread due to a rowspan associated with the back-to-back records of Isaac Westergren in Gävle, Sweden. We could ID and fix cases like this in several ways. The approach that I'm going to use here is to see if we can convert the "athlete" column into a numeric and, if so, replace these cells with the preceding value.

```
pre_iaaf <-
  pre_iaaf %>%
    mutate(athlete = ifelse(is.na(as.numeric(athlete)), athlete, lag(athlete)))
```

```
## Warning in ifelse(is.na(as.numeric(athlete)), athlete, lag(athlete)): NAs
## introduced by coercion
```

Lastly, let's fix the date column so that R recognises that the character string for what it actually is.

```
# library(lubridate) ## Already loaded
```

```
pre_iaaf <-
  pre_iaaf %>%
    mutate(date = mdy(date))
pre_iaaf
```

Finally, we have our cleaned data frame. We could easily plot the pre-IAAF data if we so wished. However, I'm going to hold off doing that until we've scraped the rest of the WR data. Speaking of which...

Your turn: Download the next two tables from the same WR100m page. Combine these two new tables with the one above into a single data frame and then plot the record progression. Answer below. (No peeking until you have tried yourself first.)

Let's start with the second table.

```
## Convert list to data_frame and clean the column names
iaaf_76 <-
  iaaf_76 %>%
  bind_rows() %>%
  as_tibble() %>%
  clean_names()
```

Fill in any missing athlete data (note that we need slightly different procedure than last time — Why?) and correct the date.

```
iaaf_76 <-
  iaaf_76 %>%
  mutate(athlete = ifelse(athlete=="", lag(athlete), athlete)) %>%
  mutate(date = mdy(date))
```

Warning: 3 failed to parse.

It looks like some dates failed to parse because a record was broken (equaled) on the same day. E.g.

```
iaaf_76 %>% tail(20)
```

```
## # A tibble: 20 x 8
##   time wind  auto athlete  nationality location_of_race date      ref
##   <dbl> <chr> <dbl> <chr>      <chr>      <chr>      <date>   <chr>
## 1  10  "2.0"  10.2 Jim Hines "United Sta~ Modesto, USA  1967-05-27 "[2]"
## 2  10  "1.8"   NA Enrique ~ "Cuba"      Budapest, Hungary 1967-06-17 "[2]"
## 3  10  "0.0"   NA Paul Nash "South Afri~ Krugersdorp, Sout~ 1968-04-02 "[2]"
## 4  10  "1.1"   NA Oliver F~ "United Sta~ Albuquerque, USA  1968-05-31 "[2]"
## 5  10  "2.0"  10.2 Oliver F~ "Charles Gr~ Sacramento, USA  1968-06-20 "[2]"
## 6  10  "2.0"  10.3 Oliver F~ "Charles Gr~ Roger Bambuck    NA      ""
## 7  9.9 "0.8"  10.0 Jim Hines "United Sta~ Sacramento, USA  1968-06-20 "[2]"
## 8  9.9 "0.9"  10.1 Ronnie R~ "United Sta~ Sacramento, USA  1968-06-20 ""
## 9  9.9 "0.9"  10.1 Charles ~ "United Sta~ Sacramento, USA  1968-06-20 ""
## 10 9.9 "0.3"  9.95 Jim Hines "United Sta~ Mexico City, Mexi~ 1968-10-14 "[2]"
## 11 9.9 "0.0"   NA Eddie Ha~ "United Sta~ Eugene, USA     1972-07-01 "[2]"
## 12 9.9 "0.0"   NA Eddie Ha~ "United Sta~ United States    NA      ""
## 13 9.9 "1.3"   NA Steve Wi~ "United Sta~ Los Angeles, USA  1974-06-21 "[2]"
## 14 9.9 "1.7"   NA Silvio L~ "Cuba"      Ostrava, Czechosl~ 1975-06-05 "[2]"
## 15 9.9 "0.0"   NA Steve Wi~ "United Sta~ Siena, Italy     1975-07-16 "[2]"
## 16 9.9 "-0.2"  NA Steve Wi~ ""          Berlin, Germany  1975-08-22 "[2]"
## 17 9.9 "0.7"   NA Steve Wi~ ""          Gainesville, USA  1976-03-27 "[2]"
## 18 9.9 "0.7"   NA Steve Wi~ "Harvey Gla~ Columbia, USA    1976-04-03 "[2]"
## 19 9.9 ""      NA Steve Wi~ ""          Baton Rouge, USA  1976-05-01 "[2]"
## 20 9.9 "1.7"   NA Don Quar~ "Jamaica"    Modesto, USA     1976-05-22 "[2]"
```

We can try to fix these cases by using the previous value. Let's test it first:

```
iaaf_76 %>%
  mutate(date = ifelse(is.na(date), lag(date), date))
```

```
## # A tibble: 54 x 8
##   time wind  auto athlete  nationality location_of_race date ref
##   <dbl> <chr> <dbl> <chr>      <chr>      <chr>      <dbl> <chr>
## 1  10.6 ""      NA Donald Lipp~ United States Stockholm, Sweden -20998 [2]
## 2  10.6 ""      NA Jackson Sch~ United States Stockholm, Sweden -18004 [2]
## 3  10.4 ""      NA Charley Pad~ United States Redlands, USA -17785 [2]
## 4  10.4 "0.0"   NA Eddie Tolan United States Stockholm, Sweden -14756 [2]
## 5  10.4 ""      NA Eddie Tolan United States Copenhagen, Denmark -14739 [2]
```



```
## 6 10.3 "" NA Percy Willi~ Canada Toronto, Ontario, ~ -14390 [2]
## 7 10.3 "0.4" 10.4 Eddie Tolan United States Los Angeles, USA -13667 [2]
## 8 10.3 "" NA Eddie Tolan Ralph Metcal~ Budapest, Hungary -13291 [2]
## 9 10.3 "" NA Eddie Tolan Eulace Peaco~ Oslo, Norway -12932 [2]
## 10 10.3 "" NA Chris Berger Netherlands Amsterdam, Netherl~ -12912 [2]
## # ... with 44 more rows
```

Whoops! Looks like all of our dates are getting converted to numbers. The reason (if you did a bit of Googling) actually has to do with the base `ifelse()` function. In this case, it's better to use the tidyverse equivalent, i.e. `if_else()`.

```
iaaf_76 <-
  iaaf_76 %>%
    mutate(date = if_else(is.na(date), lag(date), date))
iaaf_76
```

```
## # A tibble: 54 x 8
##   time wind auto athlete nationality location_of_race date ref
##   <dbl> <chr> <dbl> <chr> <chr> <chr> <date> <chr>
## 1 10.6 "" NA Donald Li~ United Stat~ Stockholm, Sweden 1912-07-06 [2]
## 2 10.6 "" NA Jackson S~ United Stat~ Stockholm, Sweden 1920-09-16 [2]
## 3 10.4 "" NA Charley P~ United Stat~ Redlands, USA 1921-04-23 [2]
## 4 10.4 "0.0" NA Eddie Tol~ United Stat~ Stockholm, Sweden 1929-08-08 [2]
## 5 10.4 "" NA Eddie Tol~ United Stat~ Copenhagen, Denma~ 1929-08-25 [2]
## 6 10.3 "" NA Percy Wil~ Canada Toronto, Ontario,~ 1930-08-09 [2]
## 7 10.3 "0.4" 10.4 Eddie Tol~ United Stat~ Los Angeles, USA 1932-08-01 [2]
## 8 10.3 "" NA Eddie Tol~ Ralph Metca~ Budapest, Hungary 1933-08-12 [2]
## 9 10.3 "" NA Eddie Tol~ Eulace Peac~ Oslo, Norway 1934-08-06 [2]
## 10 10.3 "" NA Chris Ber~ Netherlands Amsterdam, Nether~ 1934-08-26 [2]
## # ... with 44 more rows
```

Table 3: Modern Era (1977 onwards)

The final table also has its share of unique complications due to row spans, etc. You can inspect the code to see what I'm doing, but I'm just going to run through it here in a single chunk.

```
iaaf <-
  m100 %>%
    html_nodes("#mw-content-text > div > table:nth-child(19)") %>%
    html_table(fill=TRUE)

## Convert list to data_frame and clean the column names
iaaf <-
  iaaf %>%
    bind_rows() %>%
    as_tibble() %>%
    clean_names()

## Correct the date.
iaaf <-
  iaaf %>%
    mutate(date = mdy(date))

## Usain Bolt's records basically all get attributed you to Asafa Powell because
## of Wikipedia row spans (same country, etc.). E.g.
iaaf %>% tail(8)
```

```
## # A tibble: 8 x 8
##   time wind  auto athlete nationality location_of_race date      notes_note_2
##   <dbl> <chr> <dbl> <chr>    <chr>         <chr>         <date>      <chr>
## 1  9.77 1.6   9.77 Asafa ~ Jamaica Athens, Greece 2005-06-14 [2]
## 2  9.77 1.7   9.77 Justin~ United Sta~ Doha, Qatar   2006-05-12 [5][9][note~
## 3  9.77 1.5   9.76 Asafa ~ Jamaica Gateshead, Engl~ 2006-06-11 [2]
## 4  9.77 1.0   9.76 Asafa ~ 9.762 Zürich, Switzer~ 2006-08-18 [2]
## 5  9.74 1.7   9.76 Asafa ~ 9.735 Rieti, Italy   2007-09-09 [1][10]
## 6  9.72 1.7   NA    Asafa ~ Usain Bolt New York, USA 2008-05-31 [2]
## 7  9.69 0.0   9.68 Asafa ~ Asafa Powe~ Beijing, China 2008-08-16 OR[2]
## 8  9.58 0.9   9.57 Asafa ~ Asafa Powe~ Berlin, Germany 2009-08-16 CR[1][11][1~
```

Let's fix this issue

```
iaaf <-
  iaaf %>%
  mutate(
    athlete = ifelse(athlete==nationality, NA, athlete),
    athlete = ifelse(!is.na(as.numeric(nationality)), NA, athlete),
    athlete = ifelse(nationality=="Usain Bolt", nationality, athlete),
    nationality = ifelse(is.na(athlete), NA, nationality),
    nationality = ifelse(athlete==nationality, NA, nationality)
  ) %>%
  fill(athlete, nationality)
```

```
## Warning in ifelse(!is.na(as.numeric(nationality)), NA, athlete): NAs introduced
## by coercion
```

Combined eras

Let's bind all these separate eras into a single data frame. I'll use `dplyr::bind_rows()` again and select in the common variables only. I'll also add a new column describing which era an observation falls under.

```
wr100 <-
  bind_rows(
    pre_iaaf %>% select(time, athlete, nationality:date) %>% mutate(era = "Pre-IAAF"),
    iaaf_76 %>% select(time, athlete, nationality:date) %>% mutate(era = "Pre-automatic"),
    iaaf %>% select(time, athlete, nationality:date) %>% mutate(era = "Modern")
  )
wr100
```


```
## # A tibble: 99 x 7
##   time athlete nationality location_of_rac~ date      era location_of_race
##   * <dbl> <chr>    <chr>         <chr>         <date>      <chr> <chr>
## 1  10.8 Luther ~ United Sta~ Paris, France 1891-07-04 Pre~ <NA>
## 2  10.8 Cecil L~ United Kin~ Brussels, Belgi~ 1892-09-25 Pre~ <NA>
## 3  10.8 Étienne~ Belgium    Brussels, Belgi~ 1893-08-04 Pre~ <NA>
## 4  10.8 L. Atch~ United Kin~ Frankfurt/Main,~ 1895-04-13 Pre~ <NA>
## 5  10.8 Harry B~ United Kin~ Rotterdam, Neth~ 1895-08-28 Pre~ <NA>
## 6  10.8 Harald ~ Sweden    Helsingborg, Sw~ 1896-08-09 Pre~ <NA>
## 7  10.8 Isaac W~ Sweden    Gävle, Sweden   1898-09-11 Pre~ <NA>
## 8  10.8 Isaac W~ Sweden    Gävle, Sweden   1899-09-10 Pre~ <NA>
## 9  10.8 Frank J~ United Sta~ Paris, France 1900-07-14 Pre~ <NA>
## 10 10.8 Walter ~ United Sta~ Paris, France 1900-07-14 Pre~ <NA>
## # ... with 89 more rows
```

All that hard works deserves a nice plot, don't you think?

```

wr100 %>%
  ggplot(aes(x=date, y=time, col=fct_reorder2(era, date, time))) +
  geom_point(alpha = 0.7) +
  labs(
    title = "Men's 100m world record progression",
    x = "Date", y = "Time",
    caption = "Source: Wikipedia"
  ) +
  theme(legend.title = element_blank()) ## Switch off legend title

```



Summary

- Web content can be rendered either 1) server-side or 2) client-side.
- To scrape web content that is rendered server-side, we need to know the relevant CSS selectors.
- We can find these CSS selectors using SelectorGadget or, more precisely, by inspecting the element in our browser.
- We use the rvest package to read into the HTML document into R and then parse the relevant nodes.
 - A typical workflow is: `read_html(URL) %>% html_nodes(CSS_SELECTORS) %>% html_table()`.
 - You might need other functions depending on the content type (e.g. see `?html_text`).
- Just because you *can* scrape something doesn't mean you *should* (i.e. ethical and legal restrictions).
- Webscraping involves as much art as it does science. Be prepared to do a lot of experimenting and data cleaning.
- Next lecture: Webscraping: (2) Client-side and APIs.

Further resources and exercises

In the next lecture, we're going to focus on client-side web content and interacting with APIs. For the moment, you can practice your `rvest`-based scraping skills by following along with any of the many (many) tutorials available online. Lastly, we spoke a bit about the "be nice" scraping motto at the beginning of the lecture. I also wanted to point you to the **polite** package ([link](#)). It provides some helpful tools to maintain web etiquette, such as checking for permission and not hammering the host website with requests. As a bonus, it plays very nicely with the **rvest** workflow that we covered today, so please take a look.